Applied Machine Learning: HW2 (15%)

Due Monday May 3 @ 11:59pm on Canvas

Instructions: same as HW1 (same data).

1 Feature Map

1. In HW1, we binarized all categorical fields and kept the numerical fields, which makes sense for distance metrics. But for perceptron (and most other classifiers), it's actually better to binarize the numerical fields as well (so that "age" becomes many binary features such as "age=40"). Why?

2. How many binary features do you get in total after binarizing all fields? (Hint: around 230; Recall that HW1 had about 90 features). Again:

```
for i in `seq 1 9`; do cat income.train.txt.5k | cut -f $i -d ',' | sort | uniq | wc -l; done
```

Do not forget the bias dimension.

2 Perceptron and Averaged Perceptron

1. Implement the basic perceptron algorithm. Run it on the training data for 5 epochs (each epoch is a cycle over the whole training data, also known as an "iteration"). After each epoch, print the number of updates (i.e., mistakes) in this epoch (and update %, i.e., #_of_updates $/|D|$), and evaluate on the dev set and report the error rate and predicted positive rate on dev set, e.g., something like:

```
epoch 1 updates 1257 (25.1%) dev_err 23.8% (+:27.5%)
...
epoch 5 updates 1172 (23.4%) dev_err 20.5% (+:17.7%)
```

Q: what's your best error rate on dev, and at which epoch did you get it? (Hint: should be around 19%).

Actually, my dev err rates here are the upper bound. This is due to tie-breaking scenarios (when dot-product is exactly 0). Depending on how you deal with such cases (always predict +1 or always predict -1 or random), you might get 21.1% for epoch 1 and 18.7% for epoch 5 or even lower rates, but the the update rates and positive rates should be same. My dev err rates are achieved when I always consider myself wrong in tie-breaking cases (extreme unlucky case), but in reality, at testing time you don't have the correct answer in hand, and you have to make a prediction no matter what (you can't abstain even if you're uncertain), so what I did was incorrect.

On the other hand, during training, it's important to consider yourself always wrong in all tie-breaking cases; in other words, we want to be super strict to ourselves during training (e.g., study time), but allow ourselves to guess during testing (e.g., exam time).

Also note that tie-breaking is only possible with the non-averaged perceptron, as weights are all integers. With averaged perceptron, it's no longer a problem (weights are real numbers).

2. Implement the averaged perceptron. You can use either the naive version or the smart version (see slides).

   Note: the difference in speed between the two versions is not that big, as we only have around 230 features, but in HW4 the difference will be huge, where we'll have sparse features.

   Run the averaged perceptron on the training data for 5 epochs, and again report the error rate and predicted positive rate on dev set (same as above).

   Q: This time, what's your best error rate on dev, and at which epoch did you get it? (Hint: around 15%).

3. Q: What observations can you draw by comparing the per-epoch results of standard and averaged perceptrons? What about their best results?

4. Q: For the averaged perceptron, what are the five most positive/negative features? Do they make sense?

5. Q: We know that males are more likely to earn `>50K` on this dataset. But the weights for `Sex=Male` and `Sex=Female` are both negative. Why?

6. Q: What is the feature weight of the bias dimension? Does it make sense?

7. Q: Is the update % above equivalent to "training error"?

3 Comparing Perceptron with k-NN

1. What are the major advantages of perceptron over k-NN? (e.g., efficiency, accuracy, etc.)

2. Design and execute a small experiment to demonstrate your point(s).

4 Experimentations

1. Try this experiment: reorder the training data so that all positive examples come first, followed by all negative ones. Did your (vanilla and averaged) perceptron(s) degrade in terms of dev error rate(s)?

2. Try the following feature engineering:

   (a) adding the original numerical features (age, hours) as two extra, real-valued, features.
   (b) centering of each numerical dimension (**or all dimensions**) to be zero mean;
   (c) based on the above, and make each numerical dimension (**or all dimensions**) unit variance.
   (d) adding some binary combination features (e..g, edu=X-and-sector=Y)

   Q: which ones (or combinations) helped? did you get a new best error rate?

3. Collect your best model and predict on `income.test.blind` to `income.test.predicted`. The latter should be similar to the former except that the target (`>=50K`, `<50K`) field is added. Do not change the order of examples in these files.

   Q: what's your best error rate on dev? which setting achieved it? (**should be able to get 14% or less**)

   Q: what's your predicted positive % on dev? how does it compare with the ground truth positive %?

   Q: what's your predicted positive % on test?

Debriefing (required)

1. Approximately how many hours did you spend on this assignment?
2. Would you rate it as easy, moderate, or difficult?
3. Did you work on it mostly alone, or mostly with other people?
4. How deeply do you feel you understand the material it covers (0%–100%)?
5. Any other comments?